

Machine Learning Methods for Neural Data Analysis

Lecture 11: Unsupervised modeling

Announcements

- **Proposals** look great!
 - Feedback coming in the next day or two.
 - Next update due Mar 5 (next Friday).

Agenda

Finish decoding and start unsupervised modeling

- “Direct” decoders and structured prediction
- Unit III: Unsupervised models of neural and behavioral data
- Bayesian inference in latent variable models

Decoding movement from neural spike trains


A linear dynamical system (LDS) model

- One of the problems with the basic model (Lab 6 Part 1) is that it treated each time bin as independent.

- Instead, consider the following prior

$$p(X) = p(x_1) \prod_{t=2}^T p(x_t | x_{t-1})$$
$$= \mathcal{N}(x_1 | 0, Q) \prod_{t=2}^T \mathcal{N}(x_t | Ax_{t-1}, Q)$$

- Parameterized by **dynamics matrix** $A \in \mathbb{R}^{D \times D}$ and **dynamics covariance** $Q \in \mathbb{R}^{D \times D}$.



Decoding movement from neural spike trains

The posterior distribution

The posterior is given by

$$\begin{aligned}
 p(X \mid Y) &\propto \left[\mathcal{N}(x_1 \mid 0, Q) \prod_{t=2}^T \mathcal{N}(x_t \mid Ax_{t-1}, Q) \right] \left[\prod_{t=1}^T \mathcal{N}(y_t \mid Cx_t + d, R) \right] \\
 &= \exp \{ \text{"a big quadratic function of } X" \} \\
 &= \mathcal{N}(\text{vec}(X) \mid \mu, \Sigma)
 \end{aligned}$$

Where $\Sigma = J^{-1}$ and $\mu = J^{-1}h$ with,

$$J = \begin{bmatrix} J_{11} & J_{21}^\top & & & \\ J_{21} & J_{22} & J_{32}^\top & & \\ & \ddots & \ddots & \ddots & \\ & & \ddots & \ddots & J_{T,T-1}^\top \\ & & & J_{T,T-1} & J_{TT} \end{bmatrix} \quad h = \begin{bmatrix} h_1 \\ h_2 \\ \vdots \\ h_T \end{bmatrix}$$

- The blocks are given by $J_{tt} = Q^{-1} + A^\top Q^{-1}A + \mathbf{C}^\top \mathbf{R}^{-1} \mathbf{C}$ (except for J_{11} and J_{TT}), $J_{t,t-1} = -Q^{-1}A$, and $h_t = C^\top R^{-1}(y_t - d)$.

Decoding movement from neural spike trains

Poisson observations

- So far we've used a linear, Gaussian encoder for the spikes, even though they are counts!
- Suppose instead,

$$p(Y | X) = \prod_{t=1}^T \prod_{n=1}^N \text{Po} (y_{tn} | f(c_n^\top x_t + d_n))$$

- The posterior is no longer Gaussian, but it's common to approximate it as one.

Decoding movement from neural spike trains

Laplace approximation

Approximate the posterior as

$$p(X \mid Y) \approx \mathcal{N}(\mu, \Sigma)$$

where

$$\mathcal{L}(X) = -\log p(X, Y)$$

$$\mu = \operatorname{argmin}_X \mathcal{L}(X)$$

$$\Sigma = \left[\nabla^2 \mathcal{L}(X) \Big|_{X=\mu} \right]^{-1}$$

For GLM encoders, the log joint is concave and μ and Σ can be found efficiently.

Decoding movement from neural spike trains

Laplace approximation under a Poisson GLM encoder

Derive the Hessian under the Poisson GLM encoder with exponential mean function $f(a) = e^a$:

$$\begin{aligned}\frac{\partial^2}{\partial x_t \partial x_t} \mathcal{L}(X) &= J_{tt} - \sum_{n=1}^N \frac{\partial^2}{\partial x_t \partial x_t} \log \text{Po} (y_{tn} \mid f(c_n^\top x_t + d_n)) \\ &= J_{tt} - \sum_{n=1}^N \frac{\partial^2}{\partial x_t \partial x_t} \left[-f(c_n^\top x_t + d_n) + y_{tn} \log f(c_n^\top x_t + d_n) \right] \\ &= J_{tt} + \sum_{n=1}^N \exp\{c_n^\top x_t + d_n\} c_n c_n^\top\end{aligned}$$

Decoding movement from neural spike trains

Structured decoders

- If we're going to make a Gaussian approximation anyway, why not learn more flexible means and covariances?
- Recall the form of the LDS posterior,

$$J_{tt} = Q^{-1} + A^{\top} Q^{-1} A + C^{\top} R^{-1} C$$

$$J_{t,t-1} = -Q^{-1} A$$

$$h_t = C^{\top} R^{-1} (y_t - d)$$

- **Idea:** replace these with learned functions of $y_{1:T}$.

Decoding movement from neural spike trains

Structured decoders

For example,

$$p(X | Y) = \mathcal{N}(\text{vec}(X) | \mu, \Sigma)$$

$$\mu = J(Y)^{-1}h(Y)$$


$$\Sigma = J(Y)^{-1}$$

Where

$$J_{tt} = Q^{-1} + A^{\top}Q^{-1}A + f(y_{t-\Delta:t+\Delta})$$

$$J_{t,t-1} = -Q^{-1}A$$

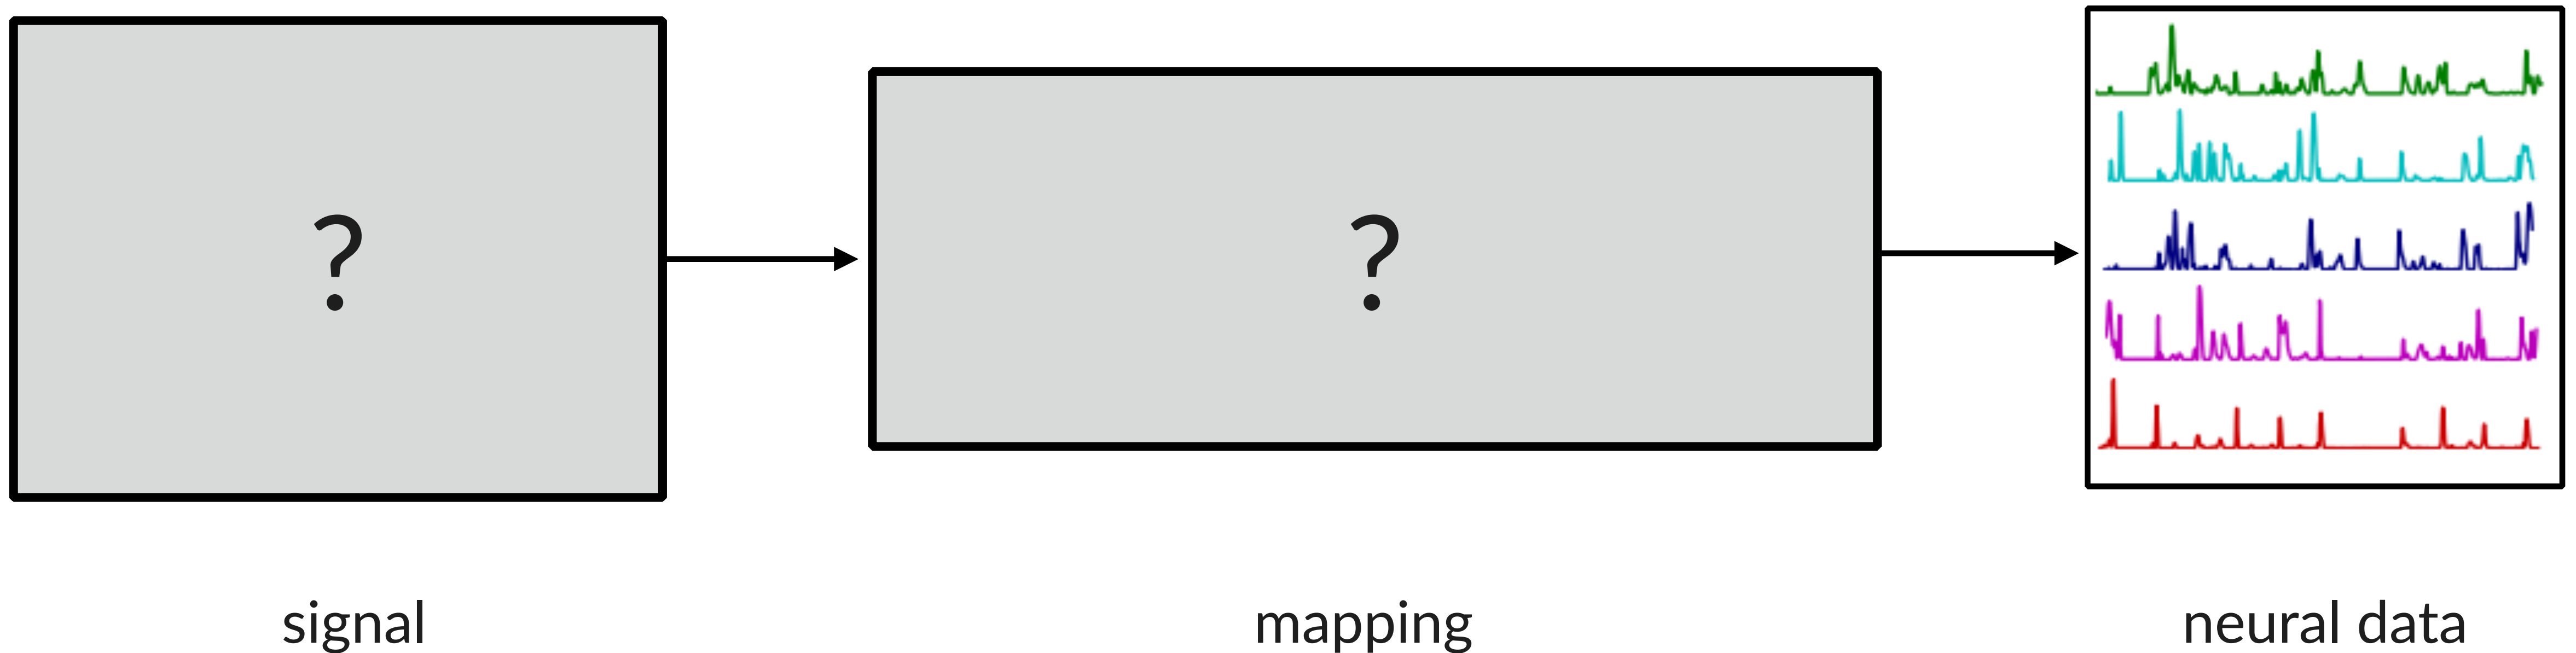
$$h_t = g(y_{t-\Delta:t+\Delta})$$



Unit III: Unsupervised modeling

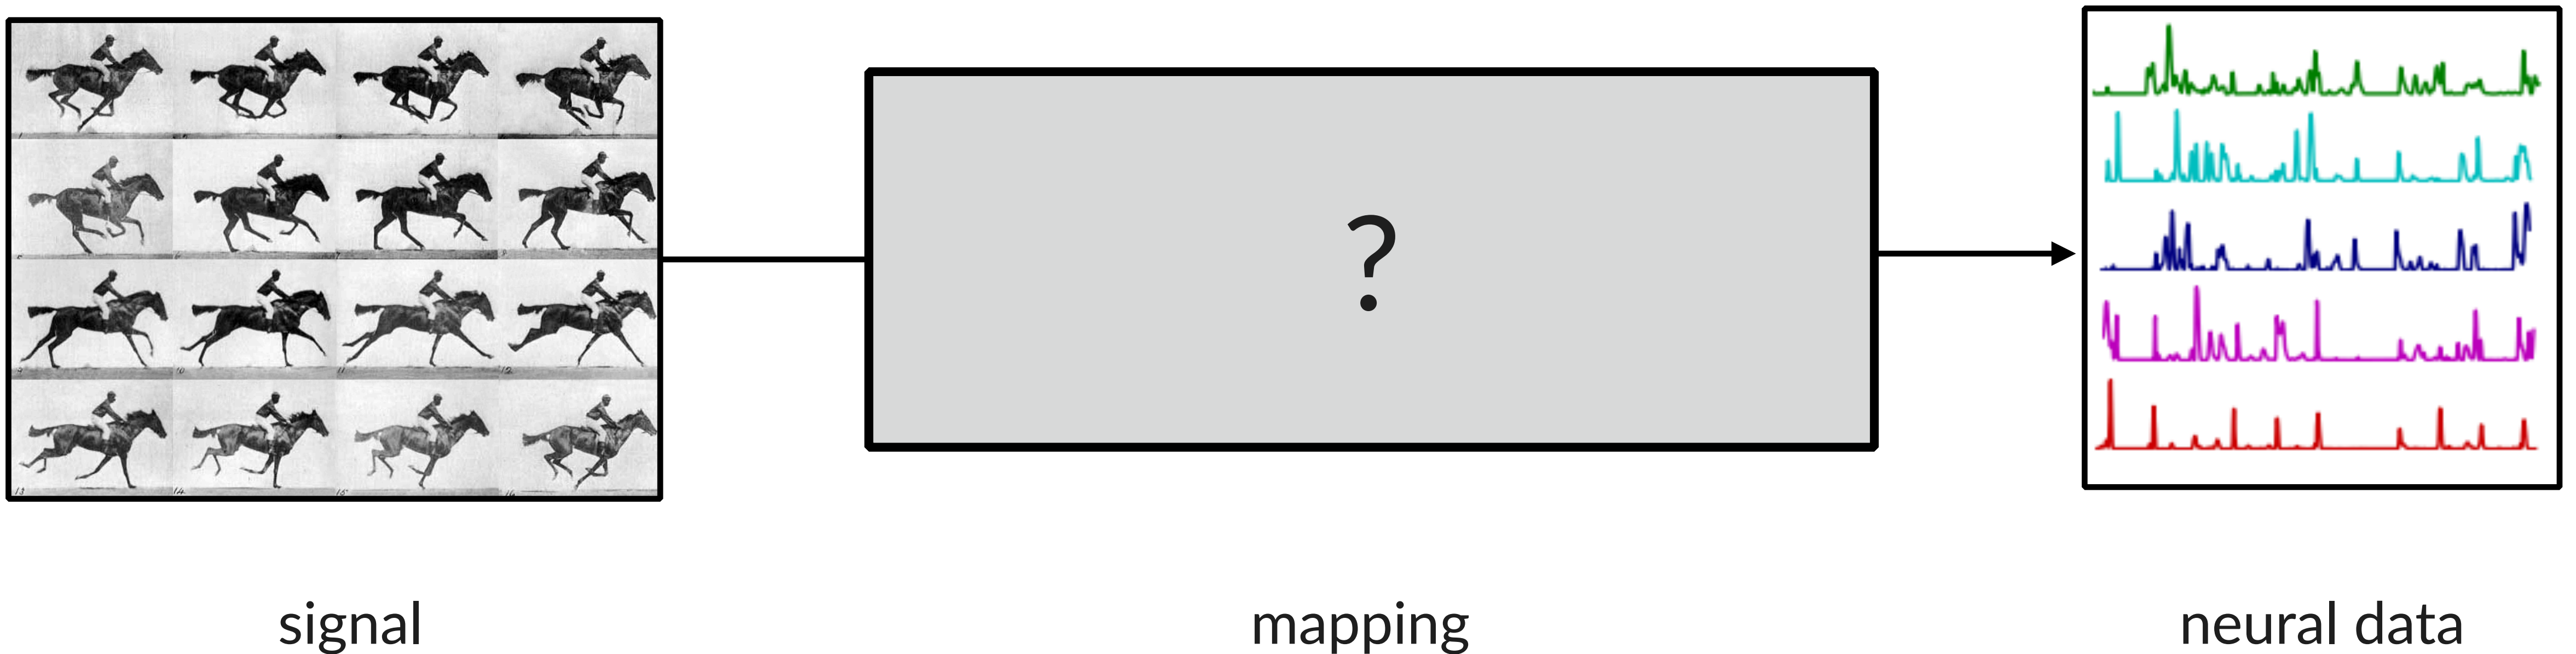
Data-driven modeling

Searching for signals to explain neural activity



Data-driven modeling

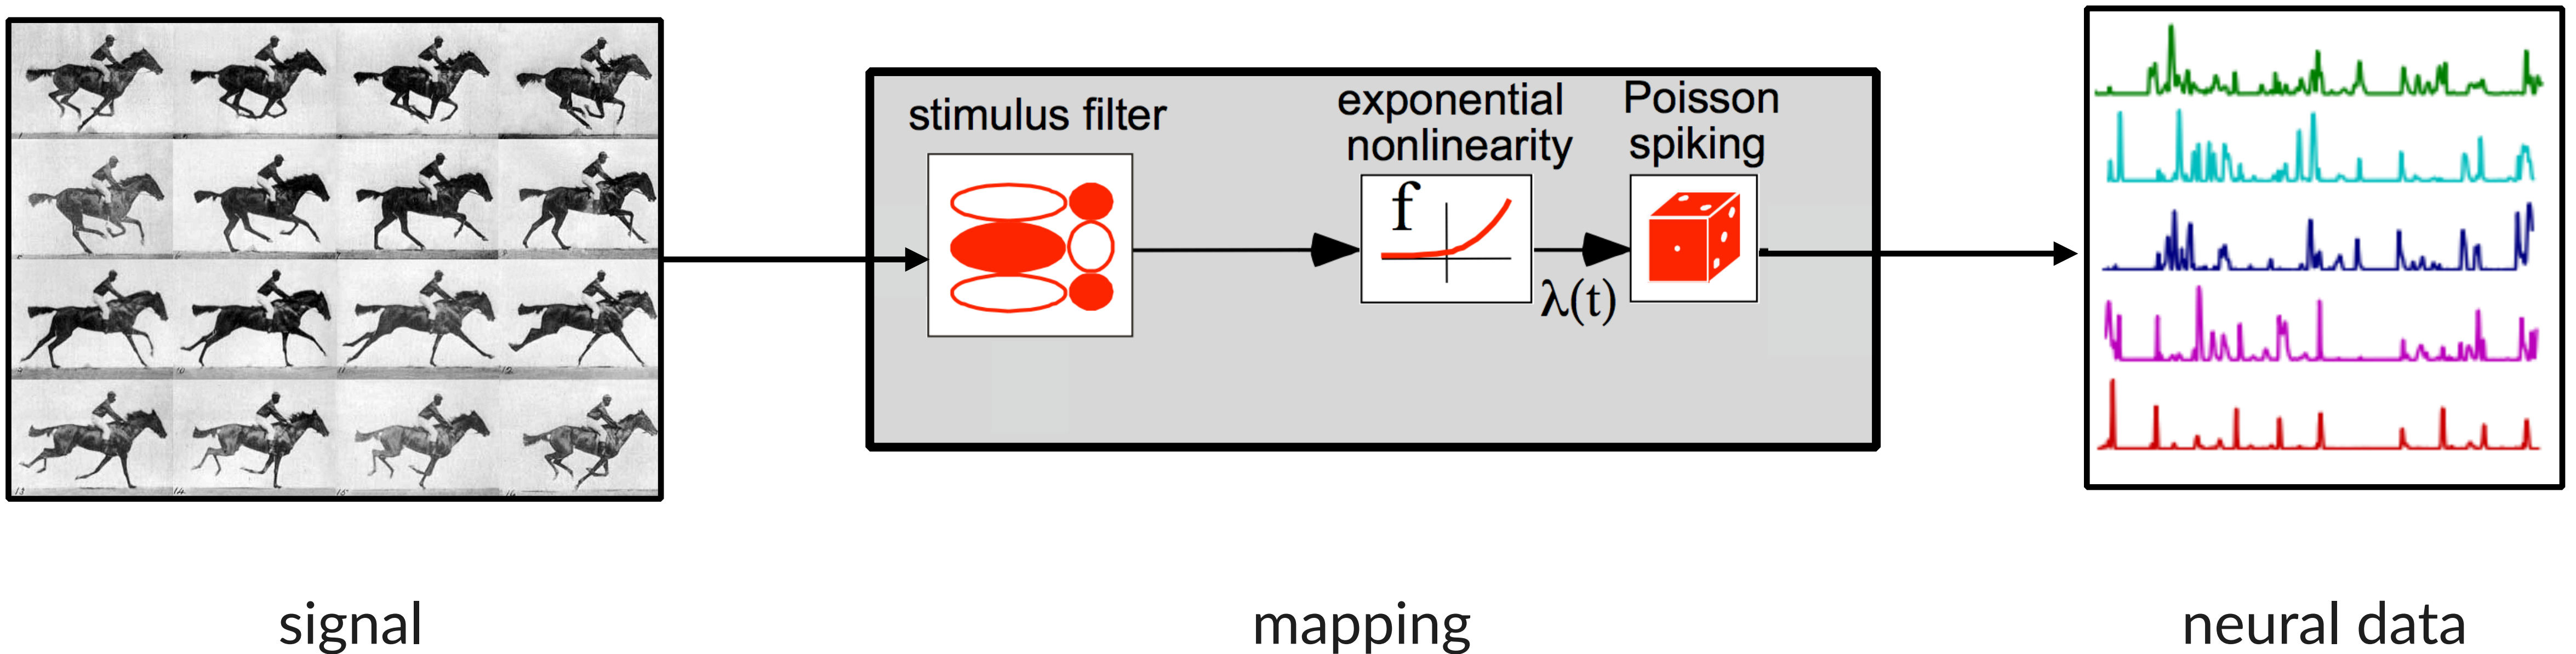
Searching for signals to explain neural activity



Encoding models: given stimulus (covariates) and response, find mapping.

Data-driven modeling

Searching for signals to explain neural activity




Recent examples: Musall et al (2018), Stringer et al (2018)


Paninski (2004)
Truccolo et al (2005)
Pillow et al (2008)

Data-driven modeling


Searching for signals to explain neural activity



signal



mapping

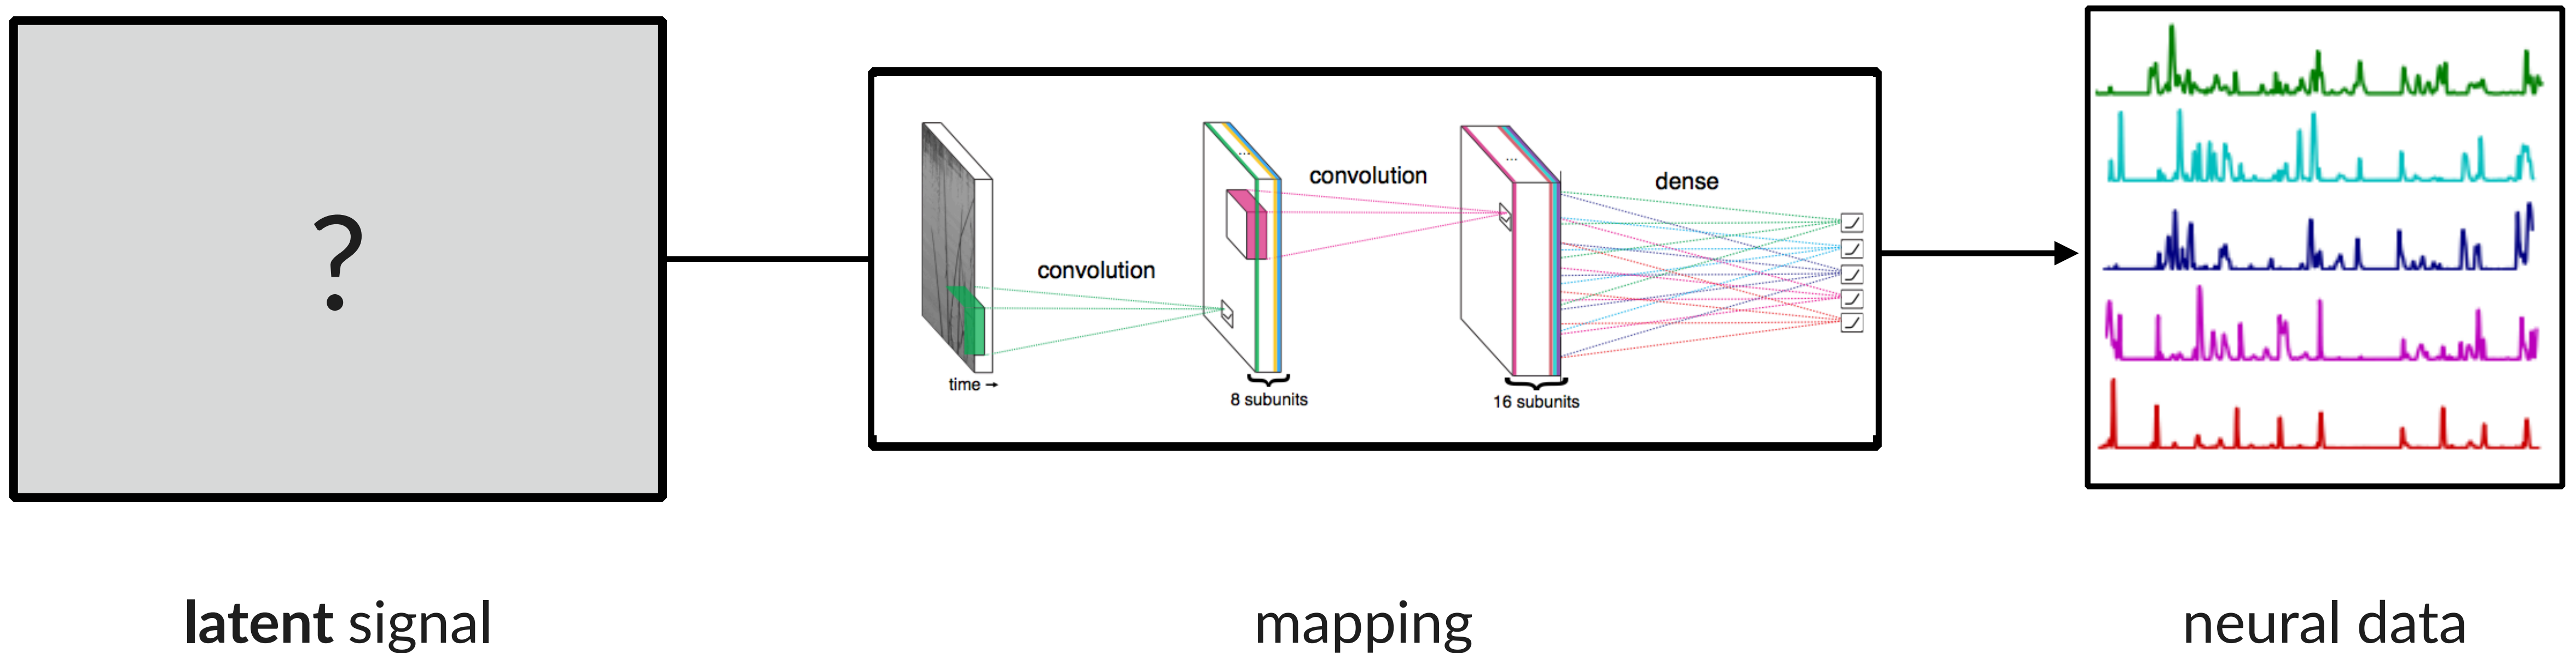


neural data

Toward nonlinear and/or more biophysically plausible mappings.

Data-driven modeling


Searching for signals to explain neural activity




Alternative: try to infer latent signals from the data

Data-driven modeling


Searching for signals to explain neural activity



latent signal



mapping



neural data

Alternative: try to infer latent signals from the data, *subject to constraints*.

Latent variable modeling is all about constraints

The five D's


- *Dimensionality*: how many latent clusters, factors, etc.?
 - *Domain*: are the latent variables discrete, continuous, bounded, sparse, etc.?
 - *Dynamics*: how do the latent variables change over time?
 - *Dependencies*: how do the latent variables relate to the observed data?
 - *Distribution*: do we have prior knowledge about the variables' probability?
-
- We've already seen some examples in Unit 1!

Latent variable modeling is all about constraints

Dynamics / Domain	Domain/Dependency/Distribution			
		Continuous Linear Gaussian	Discrete (Gen.) Linear Bernoulli/Poisson/etc.	Nonlinear Observation Models
	Discrete Markovian Categorical	HMM <i>Rabiner (1989)</i>	HMM <i>Rabiner (1989)</i>	Structured VAE <i>Johnson et al (2016)</i>
	Continuous Linear Gaussian	LDS <i>Kalman (1960)</i>	Poisson LDS <i>Smith and Brown (2003), Paninski et al (2010)</i> <i>Macke et al (2011)</i>	Deep PfLDS Archer et al (2015); Gao et al (2016)
	Continuous Nonlinear (parametric) Gaussian	NLDS, e.g. Hodgkin-Huxley <i>Ahrens, Huys, Paninski (2006)</i> <i>Huys and Paninski (2009)</i>	NLDS, e.g. Hodgkin-Huxley <i>Meng, Kramer, Eden (2011)</i>	GPSSM, DKF, LFADS, VIND <i>Frigola et al (2013)</i> , <i>Krishnan et al (2015)</i> , <i>Sussillo et al (2016)</i> , <i>Hernandez et al (2018)</i>
	Mixed Switching Linear	SLDS <i>Ghahramani and Hinton (1996)</i> <i>Murphy (1998)</i>	Poisson SLDS <i>Petreska et al (2013)</i>	Structured VAE <i>Johnson et al (2016)</i>
	Mixed Recurrent Linear	recurrent/augmented SLDS <i>Barber (2006); Pachitariu et al (2014);</i> <i>Linderman et al (2017); Nassar et al (2019)</i>	rSLDS <i>Linderman et al (2017)</i> <i>Nassar et al (2019)</i>	Structured VAE <i>Johnson et al (2016)</i>
	Continuous Nonlinear (smoothing) Gaussian	GPFA <i>Yu, Cunningham, et al (2009)</i>	vLGP <i>Zhao and Park (2017)</i>	GPLVM <i>Lawerence (2005), Wu et al (2017)</i>
	Continuous Nonlinear (nonparametric) Gaussian	GPSSM, DKF, LFADS, VIND <i>Frigola et al (2013)</i> , <i>Krishnan et al (2015)</i> , <i>Sussillo et al (2016)</i> , <i>Hernandez et al (2018)</i>	GPSSM, DKF, LFADS, VIND <i>Frigola et al (2013)</i> , <i>Krishnan et al (2015)</i> , <i>Sussillo et al (2016)</i> , <i>Hernandez et al (2018)</i>	GPSSM, DKF, LFADS, VIND <i>Frigola et al (2013)</i> , <i>Krishnan et al (2015)</i> , <i>Sussillo et al (2016)</i> , <i>Hernandez et al (2018)</i>


Motivating Example: summarizing videos with behavioral states

Frame 0




Motivating Example: summarizing videos with behavioral states


Rear down




Walk forward




Grooming




Scrunch



Rear up



Jump



Formulating as a probabilistic model

- **Variables:** Let,
 - $y_t \in \mathbb{R}^P$ denote the (vectorized) image at time t .
 - $z_t \in \{1, \dots, K\}$ denote the discrete latent state (aka behavioral “syllable”) at time t .
- **Model:** Assume each time frame is independent and,
$$z_t \sim \text{Cat}(\pi)$$
$$y_t \mid z_t \sim \mathcal{N}(d_{z_t}, R_{z_t})$$
- **Parameters:** Let $\Theta = \pi, \{d_k, R_k\}_{k=1}^K$ denote the parameters,
 - $\pi \in \Delta_K$ is the prior probability of each state
 - $(d_k, R_k) \in \mathbb{R}^P \times \mathbb{R}^{P \times P}$ are the conditional mean and variance of images for discrete state $z_t = k$.

Bayesian inference in latent variable models

Bayesian inference in latent variable models

MAP Estimation

- In Unit 1 we used *maximum a posteriori* (**MAP**) estimation to find,

$$z^{\star}, \Theta^{\star} = \arg \max_{z, \Theta} \log p(y, z, \Theta)$$

- This gave us a **point estimate** of the latent variables z and parameters Θ .
- Point estimates can lead to an **overly optimistic** view of the model.
- Specifically, MAP estimation found **the best assignment**, which may not reflect the **average performance** under the prior $p(z, \Theta)$.

Bayesian inference in latent variable models

Integrating over the latent variables

- A more **Bayesian approach** is to **integrate** over the latent variables.
- First, **learn** a point estimate of the parameters,

$$\Theta^{\star} = \arg \max_{\Theta} \log p(y, \Theta)$$

where $p(y, \Theta) = \int p(y, z, \Theta) \, dz = \mathbb{E}_{p(z, \Theta)}[p(y \mid z, \Theta)]$ is the **marginal likelihood**.

- Then, **infer** the posterior distribution over latent variables given observed data and parameters,

$$p(z \mid y, \Theta) = \frac{p(y \mid z, \Theta) p(z \mid \Theta) p(\Theta)}{p(y, \Theta)}$$

- (A “fully Bayesian” approach would integrate over both z and Θ .)

Bayesian inference in latent variable models

Maximizing the marginal likelihood

- How to learn the parameters?
- First idea: **gradient ascent**,

$$\nabla_{\Theta} \log p(y, \Theta) = \frac{\nabla_{\Theta} p(y, \Theta)}{p(y, \Theta)} = \frac{\int \nabla_{\Theta} p(y, z, \Theta) dz}{\int p(y, z, \Theta) dz}$$

- Sometimes, these integrals are available in **closed form**.
 - For example, when z **is discrete** the integrals become sums.
- Can we do better?

Bayesian inference in latent variable models

Lower bound the marginal likelihood

- Next idea: lower bound the marginal likelihood with a more tractable form,

$$\log p(y, \Theta) = \log \int p(y, z, \Theta) \mathrm{d}z$$

Bayesian inference in latent variable models

Lower bound the marginal likelihood

- Next idea: lower bound the marginal likelihood with a more tractable form,

$$\begin{aligned}\log p(y, \Theta) &= \log \int p(y, z, \Theta) \, dz \\ &= \log \int \frac{q(z)}{q(z)} p(y, z, \Theta) \, dz && \text{for any distribution } q(z) \\ &= \log \mathbb{E}_{q(z)} \left[\frac{p(y, z, \Theta)}{q(z)} \right] \\ &\geq \mathbb{E}_{q(z)} [\log p(y, z, \Theta) - \log q(z)] && \text{by Jensen's inequality} \\ &\triangleq \mathcal{L}[q, \Theta]\end{aligned}$$

- \mathcal{L} is called the **evidence lower bound** or the **ELBO** for short.

Bayesian inference in latent variable models

Coordinate ascent on the ELBO

- Update the parameters,

$$\Theta \leftarrow \arg \max_{\Theta} \mathcal{L}[q, \Theta] = \arg \max_{\Theta} \mathbb{E}_{q(z)}[\log p(y, z, \Theta)]$$

- Update the distribution on latent variables,

$$q \leftarrow \arg \max_q \mathcal{L}[q, \Theta]$$

$$= \arg \max_q \mathbb{E}_{q(z)} \left[\frac{\log p(y, z, \Theta)}{q(z)} \right]$$

$$= \arg \min_q \text{KL} \left(q(z) \parallel p(z \mid y, \Theta) \right)$$

$$= p(z \mid y, \Theta)$$

Bayesian inference in latent variable models

The Expectation-Maximization (EM) algorithm

- **M-step:** Maximize the expected log probability

$$\Theta \leftarrow \arg \max_{\Theta} \mathbb{E}_{q(z)} [\log p(y, z, \Theta)]$$


- **E-step:** Update the posterior over latent variables

$$q \leftarrow p(z \mid y, \Theta)$$

- After each E-step, the **ELBO** is tight:

$$\begin{aligned} \mathcal{L}[p(z \mid y, \Theta), \Theta] &= \mathbb{E}_{p(z|y, \Theta)} \left[\log \frac{p(y, z, \Theta)}{p(z \mid y, \Theta)} \right] \\ &= \mathbb{E}_{p(z|y, \Theta)} [\log p(y, \Theta)] \\ &= \log p(y, \Theta) \end{aligned}$$

- EM converges to **local optima** of the marginal distribution.



Bayesian inference in latent variable models

EM for the Gaussian mixture model

Recall the model,

$$z_t \sim \text{Cat}(\pi)$$
$$y_t \mid z_t \sim \mathcal{N}(d_{z_t}, R_{z_t})$$

with parameters $\Theta = \pi, \{d_k, R_k\}_{k=1}^K$.

- **E-step:** Update the posterior over latent variables,

$$q(z_t = k) \leftarrow p(z_t = k \mid y_t, \Theta) \propto \frac{\pi_k \mathcal{N}(y_t \mid d_k, R_k)}{\sum_{j=1}^K \pi_j \mathcal{N}(y_t \mid d_j, R_j)}$$

- **M-step:** Update the parameters. Let $N_k = \sum_{t=1}^T q(z_t = k)$, then

$$\pi_k \leftarrow \frac{N_k}{T}$$
$$d_k \leftarrow \frac{1}{N_k} \sum_{t=1}^T q(z_t = k) y_t$$
$$R_k \leftarrow \frac{1}{N_k} \sum_{t=1}^T q(z_t = k) (y_t - d_k)(y_t - d_k)^\top$$

Conclusion

- Unsupervised models, specifically **latent variable models**, seek simple underlying variables to explain neural or behavioral data.
- LVMs are all about **constraints**.
- MAP estimation, which we used in Unit 1, yielded a point estimate, but can be over-optimistic.
- EM maximizes the marginal likelihood by coordinate ascent on the **latent variable posterior** and the **parameters**.